Alex Rivera

AI & Machine Learning Engineer | alex.rivera@example.com | +1 (555) 019-2834 | San Francisco, CA | linkedin.com/in/alex-rivera-sample

Professional Summary

Results-driven AI/ML Engineer with 3+ years of experience designing, training, and deploying production machine learning models, retrieval-augmented generation (RAG) systems, and real-time inference APIs. Proficient in Python, PyTorch, LangChain, and MLOps tools. Adept at cross-functional collaboration to deliver robust, scalable AI architectures that drive user engagement and business efficiency.

Work Experience

AI Engineer | Nexa AI Solutions | *June 2024 – Present | Remote*

- Architected and deployed an end-to-end multi-agent orchestration pipeline using LangGraph and FastAPI, serving 30,000+ daily inference requests with sub-200ms latency.
- Implemented hybrid RAG retrieval combining dense vector embeddings and BM25 lexical search, increasing context accuracy by 28%.
- Integrated model versioning, automated evaluation gates, and continuous monitoring with MLflow and Docker.

Machine Learning Software Engineer | DataWave Technologies | *August 2022 – May 2024 | San Francisco, CA*

- Developed and maintained microservices in Python (FastAPI, Redis, PostgreSQL) supporting distributed NLP feature extraction.
- Optimized deep learning inference with ONNX Runtime and TensorRT, reducing cloud GPU hosting expenses by 35%.
- Built automated data validation and ingestion pipelines processing over 500,000 text documents daily.

Core Technical Skills

- **Languages**: Python, SQL, TypeScript, Bash
- **Frameworks & Libraries**: PyTorch, LangChain, LangGraph, Hugging Face Transformers, Scikit-learn, FastAPI, Flask
- **Data & Databases**: PostgreSQL, SQLite, MongoDB, Vector Databases (ChromaDB, Pinecone)
- **MLOps & Tools**: Docker, Kubernetes, Git, CI/CD, MLflow, Linux

Education

B.S. in Computer Science — State University (2022)